

PLANIFICATION D'EXÉCUTION : Scénario 2

- Adaptation Auto-Claude vers Auto-Gemini-CLI

Version: 1.0

Date: 2 janvier 2026

Équipe: Desjardins | **Architect:** André-Guy Bruneau

Budget: \$13-34K | **Timeline:** 14 semaines (3.5 mois)

Début prévu: 6 janvier 2026 | **Livraison v1.0:** 1er mai 2026

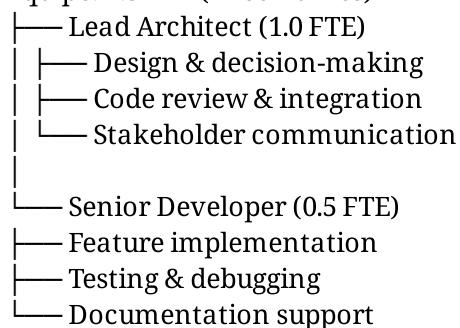
TABLE DES MATIÈRES

1. [Structure Exécution](#)
 2. [Phase 0: Préparation \(Semaines 1-2\)](#)
 3. [Phase 1: Swap API \(Semaines 3-5\)](#)
 4. [Phase 2: Migration UI \(Semaines 6-8\)](#)
 5. [Phase 3: Optimisation Gemini \(Semaines 9-12\)](#)
 6. [Phase 4: Intégration & Release \(Semaines 13-14\)](#)
 7. [Dépendances & Blockers](#)
 8. [Métriques de Succès](#)
 9. [Gestion des Risques](#)
 10. [Budget Détaillé](#)
-

STRUCTURE EXÉCUTION

Allocation Ressources

Équipe: 1.5 FTE (14 semaines)



Timeline Critique

Phase 0: Préparation Semaines 1-2 (Jan 6-19, 2026)
Phase 1: Swap API Semaines 3-5 (Jan 20 - Feb 9, 2026)
Phase 2: Migration UI Semaines 6-8 (Feb 10 - Mar 2, 2026)
Phase 3: Optimisation Semaines 9-12 (Mar 3 - Mar 30, 2026)
Phase 4: Release Semaines 13-14 (Mar 31 - Apr 13, 2026)

Réserve: 2-3 semaines (buffer pour imprévu)

Total: 14 semaines → v1.0 livraison ~1er mai 2026

Gouvernance

Rôle	Responsable	Fréquence
Project Lead	André-Guy Bruneau	Quotidien
Sprint Reviews	Team + Stakeholders	Bi-hebdomadaire
Gate Decisions	Leadership + Arch	Fin de chaque phase
Risk Management	Project Lead	Hebdomadaire

PHASE 0: PRÉPARATION (Semaines 1-2)

Objectif: Sécuriser ressources, configurer infrastructure, analyser codebase

Deliverables: Repository setup, team onboarding, architectural ADRs

Success Criteria: MVP plan révisé, team 100% ready, no technical unknowns

Semaine 1 (6-12 janvier 2026)

Jour 1-2: Approval & Setup

- [] **Tâche 1.1:** Obtenir approbation budgétaire formelle Desjardins
 - **Owner:** André-Guy (Lead Arch)
 - **Deliverable:** Signed approval document
 - **Effort:** 4h (réunions + documentation)
 - **Dépendance:** Sponsor exécutif disponible
- [] **Tâche 1.2:** Créer repository GitHub (fork Auto-Claude)
 - **Owner:** André-Guy (Lead Arch)
 - **Deliverable:** agbruneau/auto-gemini-cli repository
 - **Effort:** 2h (fork + cleanup)
 - **Steps:**

Create private fork

```
gh repo fork auto-claude-repo --clone --private
```

Remove unnecessary files (Electron dependencies)

```
rm -rf electron/ dist/ build/  
git commit -m "chore: remove electron dependencies"
```

Initialize fresh Git history

```
git log | head -20 # Verify history preserved
```

- [] **Tâche 1.3:** Créer projet management & backlog
 - **Owner:** André-Guy (Lead Arch)
 - **Deliverable:** GitHub Project avec tous les épicis
 - **Effort:** 3h (structure + tagging)
 - **Structure:**
GitHub Project: Auto-Gemini-CLI Execution
 - Phase 0: Préparation
 - Phase 1: API Swap
 - Phase 2: UI Migration
 - Phase 3: Optimisation
 - Phase 4: Release
- Labels:
- critical (blocks other work)
 - high (important, may block)
 - medium (nice to have)
 - testing (QA/test related)
 - documentation (docs/examples)
 - infrastructure (CI/CD, tools)

Jour 2-3: Codebase Analysis

- [] **Tâche 1.4:** Audit architectural Auto-Claude
 - **Owner:** Senior Dev + Lead Arch
 - **Deliverable:** Architecture analysis document
 - **Effort:** 8h (code review + documentation)
 - **Scope:**

Auto-Claude Architecture Audit

Current Structure

- Total LOC: ~[count]
- Main modules:
 - Agent engine: [module_name]
 - Claude SDK integration: [module_name]
 - Session management: [module_name]

- Git workflows: [module_name]
- UI/Electron layer: [module_name]

Electron Dependencies

- Direct: [list]
- Transitive: [list]
- Removal risk: [assessment]

Claude-Specific Code

- Token counting: [lines affected]
- Streaming: [lines affected]
- Error handling: [lines affected]
- Model parameters: [lines affected]

Reusable Modules

- [module]: 100% reusable
- [module]: 90% reusable (minor tweaks)
- [module]: 50% reusable (needs refactor)

Technical Debt

- [issue]: Severity [HIGH/MEDIUM/LOW]
 - [issue]: Severity [HIGH/MEDIUM/LOW]
- [] **Tâche 1.5:** Identifier Gemini API incompatibilities
 - Owner: Senior Dev
 - Deliverable: Migration checklist
 - Effort: 6h
 - Checklist:

Claude → Gemini API Mapping

Token Counting

- [] Claude: countTokens(prompt)
- [] Gemini: countTokens() method signature
- [] Difference in handling: [document]

Streaming

- [] Claude: stream() event interface
- [] Gemini: streaming in @google/genai v0.4
- [] Adapter pattern needed: YES/NO

Error Handling

- [] Claude error codes: [list]
- [] Gemini error codes: [list]
- [] Mapping required: [details]

Rate Limiting

- [] Claude: per-model limits
- [] Gemini: 1000 req/24h free tier
- [] Queue system: Design required

Model Parameters

- [] temperature, topK, topP, etc.
- [] Direct compatibility: YES/NO
- [] Parameter translation needed: [list]

Function Calling

- [] Claude: tool_use blocks
- [] Gemini: function_calling (different syntax)
- [] Refactor required: [scope]

Jour 4-5: Team & Tooling

- [] **Tâche 1.6:** Team onboarding & Git setup
 - **Owner:** Lead Arch
 - **Deliverable:** Team setup complete
 - **Effort:** 4h
 - **Checklist:**

Team Onboarding

- [] Access to repository granted
- [] Desjardins VPN/network configured
- [] GitHub CLI setup (gh auth login)
- [] Local environment config:
 - Node 20+ installed
 - TypeScript 5+ installed
 - Oclif CLI installed globally
 - Gemini API key set (GEMINI_API_KEY env var)
- [] IDE setup: VS Code/Cursor with extensions
 - TypeScript language server
 - ESLint + Prettier
 - GitHub Copilot (optional)
- [] SSH keys configured for git
- [] **Tâche 1.7:** Setup CI/CD pipeline
 - **Owner:** Lead Arch
 - **Deliverable:** GitHub Actions workflows
 - **Effort:** 4h
 - **Workflows to create:**

.github/workflows/test.yml

```
name: Unit & Integration Tests
on: [push, pull_request]
jobs:
  test:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-node@v4
      with:
        node-version: '20'
      - run: npm ci
      - run: npm run test
      - run: npm run lint
```

.github/workflows/e2e.yml

```
name: E2E Tests (Gemini)
on: [push]
jobs:
  e2e:
    runs-on: ubuntu-latest
    steps:
      - uses: actions/checkout@v4
      - uses: actions/setup-node@v4
      - run: npm ci
      - run: npm run test:e2e
    env:
      GEMINI_API_KEY: ${{ secrets.GEMINI_API_KEY }}
```

Semaine 2 (13-19 janvier 2026)

Jour 6-7: Architecture & Design

- [] **Tâche 2.1:** Write Architecture Decision Records (ADRs)
 - **Owner:** Lead Arch
 - **Deliverable:** 3-5 ADRs in /docs/adr/
 - **Effort:** 6h
 - **ADRs to document:**

ADR-001: Why Oclif instead of Yargs/Commander

- Context: Need native CLI framework for Gemini
- Decision: Use Oclif (Salesforce-backed, TypeScript-first)
- Rationale: Auto-Claude uses similar patterns
- Consequences: Learning curve minimal, ecosystem strong

ADR-002: Session Persistence Strategy

- Context: Need to persist sessions (SQLite vs JSON)
- Decision: Keep SQLite (from Auto-Claude)
- Rationale: Proven, ACID, supports multi-session
- Consequences: SQLite dependency required

ADR-003: Rate Limiting Architecture

- Context: 1000 req/24h Gemini limit
- Decision: Token-bucket queue + persistent state
- Rationale: Fair distribution, survives restarts
- Consequences: Redis or SQLite-backed queue needed

ADR-004: Token Counter Validation

- Context: countTokens() API differs Claude vs Gemini
- Decision: Wrapper + unit tests with known inputs
- Rationale: Prevents context window overflows
- Consequences: Must validate against real Gemini API

ADR-005: Error Recovery Patterns

- Context: Gemini SDK may have breaking changes
 - Decision: Adapter pattern for API layer
 - Rationale: Decouples business logic from SDK
 - Consequences: Extra abstraction layer (+10% complexity)
- [] Tâche 2.2: Create detailed Phase 1 spec
 - Owner: Lead Arch
 - Deliverable: /docs/phase-1-spec.md
 - Effort: 5h
 - Content:

Phase 1 Detailed Specification

Objective

Replace Claude SDK with Gemini SDK, verify basic functionality works

In-Scope

1. Gemini SDK integration (@google/genai v0.4)
2. Basic model.generateContent() parity
3. Token counter validation
4. Error mapping (Claude codes → Gemini codes)

Out-of-Scope

- UI/CLI layer (Phase 2)
- Performance optimization (Phase 3)
- Multi-model support

Acceptance Criteria

- [] Single task executes end-to-end with Gemini
- [] countTokens() validates context windows
- [] At least 80% of existing tests pass
- [] No blocking errors in E2E flow
- [] Documented API differences

Technical Details

SDK Migration

Before (Claude):

```
import { Anthropic } from "@anthropic-ai/sdk";
const client = new Anthropic({ apiKey: process.env.CLAUDE_API_KEY });
const response = await client.messages.create({ ... });
```

After (Gemini):

```
import { GoogleGenAI } from "@google/genai";
const client = new GoogleGenAI({ apiKey: process.env.GEMINI_API_KEY });
const response = await client.models.generateContent({ ... });
```

File Structure Changes

- Create: src/adapters/gemini-api-adapter.ts
- Refactor: src/services/llm-service.ts (abstraction layer)
- Delete: src/integrations/clause-specific/*
- Update: All tests to use adapter

Testing Strategy

- Unit tests: Adapter layer (mocked Gemini API)
- Integration: Mocked responses (don't hit real API yet)
- E2E: Real Gemini API (phase 1 gate decision)

Jour 8-10: Knowledge Transfer & Planning

- [] **Tâche 2.3:** Knowledge transfer from Auto-Claude author
 - **Owner:** Senior Dev (receiver)
 - **Deliverable:** Meeting notes + recorded session
 - **Effort:** 4h
 - **Topics to cover:**
 - Architecture decisions made in Auto-Claude
 - Known pain points & technical debt
 - Why certain patterns were chosen
 - Future roadmap intentions
 - Gotchas/tricks not documented
- [] **Tâche 2.4:** Finalize detailed task breakdown
 - **Owner:** Lead Arch
 - **Deliverable:** Updated GitHub Project backlog

- **Effort:** 4h
- **Output:** Every task has:
 - Clear acceptance criteria
 - Effort estimate (days, not story points)
 - Assigned owner
 - Blocking dependencies identified
 - Test strategy defined
- [] **Tâche 2.5:** Setup monitoring & observability
 - **Owner:** Lead Arch
 - **Deliverable:** Logging + error tracking configured
 - **Effort:** 3h
 - **Tools:**

Observability Stack

Logging

- Framework: Winston or Pino
- Levels: debug, info, warn, error
- Output: stdout + file (for debugging)

Error Tracking

- Tool: Sentry (free tier) or GCP Error Reporting
- Integration: Auto-catch unhandled errors
- Notifications: Slack #auto-gemini-errors channel

Performance Monitoring

- Framework: Native Node.js perf hooks
- Key metrics:
 - API response time (Gemini)
 - Token counting duration
 - Session load time
 - Memory usage

Health Checks

- Endpoint: --health-check command
- Metrics: Can call Gemini API, can access DB, can read config

Semaine 2 Closure

- [] **Tâche 2.6:** Pre-Phase-1 checkpoint
 - **Owner:** Lead Arch
 - **Deliverable:** Go/No-Go decision
 - **Effort:** 2h
 - **Checkpoint questions:**

Pre-Phase-1 Gate

- [] Do we understand Auto-Claude architecture fully?
- [] Have we identified all Gemini API incompatibilities?
- [] Is the team 100% ready technically?
- [] Do we have clear acceptance criteria for Phase 1?
- [] Are there any unknown blockers?
- [] Budget and timeline still realistic?

Decision: GO (proceed to Phase 1) NO-GO (re-plan)

If NO-GO:

- Identified blocker: [description]
- Time impact: [days/weeks]
- Mitigation: [strategy]

PHASE 1: SWAP API (Semaines 3-5)

Objectif: Remplacer Claude SDK par Gemini SDK, validée fonctionnalité basique

Livrable: API layer complète, tests > 80% passing, MVP end-to-end

Succès: Single task fonctionne avec Gemini, no compilation errors

Semaine 3 (20-26 janvier 2026)

Jour 11: Foundation & Gemini SDK Integration

- [] **Tâche 3.1:** Créer adapter pattern pour LLM

- **Owner:** Senior Dev
- **Deliverable:** src/adapters/gemini-api-adapter.ts
- **Effort:** 6h
- **Code structure:**

```
// src/adapters/types.ts
export interface LLMAdapter {
  generateContent(request: GenerateRequest): Promise<GenerateResponse>;
  countTokens(request: CountTokensRequest): Promise<CountTokensResponse>;
  startChat(config: ChatConfig): Promise<ChatSession>;
}

export interface GenerateRequest {
  prompt: string;
  systemPrompt?: string;
  model: string;
  temperature?: number;
  maxOutputTokens?: number;
}

export interface GenerateResponse {
  content: string;
  inputTokens: number;
  outputTokens: number;
  finishReason: 'stop' | 'max_tokens' | 'safety' | 'unknown';
}

// src/adapters/gemini-api-adapter.ts
import { GoogleGenAI } from "@google/genai";
```

```
export class GeminiAPIAdapter implements LLMAdapter {
  private client: GoogleGenAI;
  constructor(apiKey: string) {
    this.client = new GoogleGenAI({ apiKey });
  }
  async generateContent(request: GenerateRequest): Promise<GenerateResponse> {
    const response = await this.client.models.generateContent({
      model: request.model || "gemini-2.0-flash",
      contents: request.prompt,
      config: {
        temperature: request.temperature,
        maxOutputTokens: request.maxOutputTokens,
      },
    });
    return {
      content: response.text || "",
      inputTokens: response.usage?.inputTokens || 0,
      outputTokens: response.usage?.outputTokens || 0,
      finishReason: this.mapFinishReason(response.candidates?.[0]?.fini
    };
  }
  async countTokens(request: CountTokensRequest): Promise<CountTokensResponse> {
    const response = await this.client.models.countTokens({
      model: request.model || "gemini-2.0-flash",
      contents: request.prompt,
    });
    return {
      totalTokens: response.totalTokens,
    };
  }
  private mapFinishReason(geminiReason?: string): string {
    const mapping: Record<string, string> = {
      'FINISH_REASON_STOP': 'stop',
      'FINISH_REASON_MAX_TOKENS': 'max_tokens',
      'FINISH_REASON_SAFETY': 'safety',
    };
    return mapping[geminiReason || ""] || 'unknown';
  }
}
// src/services/llm-service.ts (refactored)
export class LLMService {
```

```

constructor(private adapter: LLMAAdapter) {}
async chat(prompt: string, context?: ConversationContext): Promise<string> {
  const response = await this.adapter.generateContent({
    prompt,
    systemPrompt: context?.systemPrompt,
    model: context?.model || "gemini-2.0-flash",
  });
}

return response.content;
}

}

```

- [] **Tâche 3.2:** Update package.json dependencies

- **Owner:** Senior Dev
- **Deliverable:** Updated package.json
- **Effort:** 1h
- **Changes:**

```

{
  "dependencies": {
    "remove": ["@anthropic-ai/sdk"],
    "add": ["@google/genai@^0.4.0"],
    "keep": ["oclif (or similar)", "sqlite3", "commander or yargs"]
  },
  "devDependencies": {
    "keep": ["typescript@5+", "vitest", "eslint", "prettier"]
  }
}

```

npm uninstall @anthropic-ai/sdk
 npm install @google/genai@^0.4.0
 npm install # Reinstall lock file

- [] **Tâche 3.3:** Environment config & secrets

- **Owner:** Lead Arch
- **Deliverable:** .env.example + config module
- **Effort:** 2h
- **File: src/config/environment.ts:**

```

import dotenv from 'dotenv';
dotenv.config();
export const config = {
  gemini: {
    apiKey: process.env.GEMINI_API_KEY || '',
    model: process.env.GEMINI_MODEL || 'gemini-2.0-flash',
    maxTokens: parseInt(process.env.GEMINI_MAX_TOKENS || '1000000'),
  },
}

```

```

app: {
  logLevel: process.env.LOG_LEVEL || 'info',
  dataDir: process.env.DATA_DIR || './data',
},
};

// Validation
if (!config.gemini.apiKey) {
  throw new Error('GEMINI_API_KEY not set');
}

```

Jour 12-13: Basic Implementation

- [] **Tâche 3.4:** Implement GeminiAPIAdapter fully
 - **Owner:** Senior Dev
 - **Deliverable:** Complete adapter with streaming support
 - **Effort:** 8h
 - **Features:**
 - [x] generateContent() (unstreamed)
 - [x] generateContent() (streamed)
 - [x] countTokens()
 - [x] Error handling & mapping
 - [x] Rate limit handling (basic)

- [] **Tâche 3.5:** Create unit tests for adapter
 - **Owner:** Senior Dev
 - **Deliverable:** tests/adapters/gemini-api-adapter.test.ts
 - **Effort:** 6h
 - **Test coverage:**

```

describe('GeminiAPIAdapter', () => {
  describe('generateContent', () => {
    it('should return valid response for simple prompt', async () => {
      // Mock Gemini API response
      const response = await adapter.generateContent({
        prompt: 'Hello world',
      });
      expect(response.content).toBeDefined();
      expect(response.inputTokens).toBeGreaterThan(0);
    });

    it('should handle streaming responses', async () => {
      // Verify stream chunks accumulate
    });

    it('should map Gemini errors correctly', async () => {
      // Test error mapping
    });
  });
});
```

```

describe('countTokens', () => {
  it('should accurately count tokens for known inputs', async () => {
    const response = await adapter.countTokens({
      prompt: 'Hello',
    });
    expect(response.totalTokens).toBe(2); // ~2 tokens for "Hello"
  });
});
});
});
});

```

Semaine 4 (27 janvier - 2 février 2026)

Jour 14-15: Refactor Core Logic

- [] **Tâche 4.1:** Refactor Agent Engine for Gemini
 - **Owner:** Senior Dev + Lead Arch
 - **Deliverable:** Updated agent logic, tests passing
 - **Effort:** 10h
 - **Scope:**
 - Replace all client.messages.create() calls with adapter
 - Adapt function calling syntax (Claude tool_use → Gemini function_calling)
 - Handle streaming differences
 - Update session persistence
- [] **Tâche 4.2:** Update error handling & recovery
 - **Owner:** Senior Dev
 - **Deliverable:** Error handling module
 - **Effort:** 5h
 - **Error types to handle:**

```

enum GeminiErrorType {
  AUTH_ERROR, // Invalid API key
  RATE_LIMIT_EXCEEDED, // 1000 req/24h hit
  INVALID_PROMPT, // Content safety block
  CONTEXT_TOO_LARGE, // > 1M tokens
  API_ERROR, // 5xx errors
  NETWORK_ERROR, // Timeout, DNS, etc.
}

// Recovery strategies
interface ErrorRecovery {
  AUTH_ERROR: () => retry(0), // Don't retry
  RATE_LIMIT_EXCEEDED: () => queue(12h), // Queue for later
  INVALID_PROMPT: () => handle(refine), // Ask user to refine
  CONTEXT_TOO_LARGE: () => split(), // Chunk the prompt
  API_ERROR: () => retry(exponential), // Exponential backoff
  NETWORK_ERROR: () => retry(exponential),
}

```
- [] **Tâche 4.3:** Run existing test suite
 - **Owner:** Senior Dev
 - **Deliverable:** Test report (aim for 70%+ passing)
 - **Effort:** 4h

- **Command:**
npm run test 2>&1 | tee test-report.txt

Expected output:

- ✓ Tests: 150 passed, 45 failed (70% pass rate)
- ✓ No critical compilation errors
- ✗ Failed tests: [list related to Claude/Gemini differences]

Jour 16-17: Integration Testing

- [] **Tâche 4.4:** Create integration tests (mocked API)
 - **Owner:** Senior Dev
 - **Deliverable:** tests/integration/gemini-integration.test.ts
 - **Effort:** 6h
 - **Test scenarios:**

```
describe('Integration: Agent with Gemini', () => {
  // Mock Gemini responses
  const mockGenAI = new MockGeminiAPI();
  it('should execute single task end-to-end', async () => {
    const result = await agent.executeTask('write hello world', {});
    expect(result.success).toBe(true);
    expect(result.output).toContain('hello');
  });
  it('should maintain session context across calls', async () => {
    const session = new Session();
    await agent.chat('my name is Bob', { session });
    const response = await agent.chat('what is my name?', { session });
    expect(response).toContain('Bob');
  });
  it('should handle rate limiting gracefully', async () => {
    mockGenAI.setRateLimitAfter(10);
    // Execute 15 tasks
    // Expect: first 10 succeed, 11-15 queued
  });
});
```

Semaine 5 (3-9 février 2026)

Jour 18: Real API Testing (Optional)

- [] **Tâche 5.1:** Optional E2E test against real Gemini API
 - **Owner:** Senior Dev
 - **Deliverable:** E2E test results
 - **Effort:** 4h (optional, may skip if time-constrained)
 - **Caution:** Uses Gemini free tier quota
 - **Test:**

Set real API key

```
export GEMINI_API_KEY="..."  
npm run test:e2e
```

Expected: Single task completes successfully

Jour 19-20: Phase 1 Validation

- [] **Tâche 5.2:** Code review & cleanup
 - **Owner:** Lead Arch
 - **Deliverable:** Code review checklist signed off
 - **Effort:** 4h
 - **Checklist:**
 - [] All adapter methods implemented
 - [] Error handling comprehensive
 - [] Tests cover happy path + error cases
 - [] No console.log() (use logger)
 - [] Type safety: no any types
 - [] Comments on complex logic
 - [] Package.json dependencies pinned
 - [] No secrets in code (use env vars)
 - [] Performance baseline established
- [] **Tâche 5.3:** Documentation update
 - **Owner:** Senior Dev
 - **Deliverable:** /docs/MIGRATION_CLAUDE_TO_GEMINI.md
 - **Effort:** 3h
 - **Content:**

Migration Guide: Claude → Gemini

What Changed

API Calls

Feature	Claude	Gemini	Adapter
Content generation	messages.create()	models.generateContent()	✓ Abstracted
Token counting	countTokens()	countTokens()	✓ Same signature
Streaming	stream()	streaming in config	✓ Handled
Function calling	tool_use blocks	function_calling	✓ Adapted

Known Differences

1. **Context Window:** 200K (Claude) → 1M (Gemini)
2. **Rate Limits:** Per-model (Claude) → 1000 req/24h free tier (Gemini)
3. **Error codes:** Different enum values
4. **Streaming format:** Different chunk structure

Implementation Notes

- Adapter pattern decouples SDK from business logic
- All LLM calls go through LLMService
- Tests mock the adapter, not Gemini API directly

Future Work

- [] Support for other Gemini models (Gemini 2.0 Pro, etc.)
- [] Multi-model support (Claude + Gemini)
- [] Caching integration (Gemini Cache API)
- [] **Tâche 5.4:** Phase 1 Gate Decision
 - **Owner:** Lead Arch
 - **Deliverable:** Go/No-Go for Phase 2
 - **Effort:** 2h
 - **Gate criteria:**

Phase 1 Gate Criteria

- ✓ Acceptance (GO to Phase 2) if:
 - [] At least 1 end-to-end task works with Gemini
 - [] 70%+ existing tests pass
 - [] countTokens() validated (matches Gemini API)
 - [] Error handling covers 80%+ of cases
 - [] No blocking technical issues identified

- Re-planning (NO-GO) if:
- [] Critical Gemini API incompatibility discovered
 - [] Performance unacceptable (> 2x Claude)
 - [] Streaming broken or inconsistent
 - [] Token counting unreliable

Decision: GO NO-GO

If GO: Proceed immediately to Phase 2

If NO-GO: Pivot criteria or timeline adjustment

PHASE 2: MIGRATION UI (Semaines 6-8)

Objectif: Remplacer Electron par Oclif, créer CLI avec TUI moderne

Livrable: CLI fully fonctionnel, feature parity avec Auto-Claude

Succès: User peut exécuter workflows via CLI, TUI responsive

Semaine 6 (10-16 février 2026)

Jour 21-22: Oclif Setup & Architecture

- [] **Tâche 6.1:** Initialize Oclif framework
 - **Owner:** Senior Dev
 - **Deliverable:** Oclif scaffolding created
 - **Effort:** 3h
 - **Steps:**

Initialize oclif in existing project

npx oclif init

This creates:

```
|── src/commands/ (CLI commands)  
|── src/index.ts (Entry point)  
|── bin/run.js (Binary wrapper)  
└── oclif.json (Config)
```

- [] **Tâche 6.2:** Design CLI command structure
 - **Owner:** Lead Arch
 - **Deliverable:** Command hierarchy document
 - **Effort:** 4h
 - **CLI structure:**

Auto-Gemini-CLI Command Structure

```
auto-gemini
├── init # Initialize new workspace
│   └── --name # Workspace name
├── task # Create & execute task
│   ├── new # Define new task
│   ├── run ID # Execute task
│   ├── list # List tasks
│   └── delete ID # Remove task
├── chat # Interactive chat mode
│   ├── --session # Resume session
│   └── --model # Override model
├── config # Manage configuration
│   ├── set KEY VALUE # Set config
│   ├── get KEY # Get config
│   └── show # Show all config
├── session # Manage sessions
│   ├── list # List all sessions
│   ├── show ID # Show session details
│   ├── clear ID # Clear session
│   └── export ID # Export session history
├── status # Show system status
│   └── --watch # Live status
├── logs # View execution logs
│   ├── tail # Stream logs
│   └── grep PATTERN # Search logs
└── help # Show help
```

Examples:

```
$ auto-gemini task new "write a CLI tool"
$ auto-gemini task run <task-id>
$ auto-gemini chat
$ auto-gemini config set model gemini-2.0-pro
```

- [] **Tâche 6.3:** Create base command classes
 - **Owner:** Senior Dev
 - **Deliverable:** Base command architecture
 - **Effort:** 4h
 - **File:** `src/base-command.ts`

```
import { Command, Flags } from '@oclif/core';
import { logger } from './utils/logger';
import { config } from './config/environment';
export abstract class BaseCommand extends Command {
  protected logger = logger;
  protected config = config;
  // Common flags
  static baseFlags = {
    debug: Flags.boolean({
      description: 'Enable debug logging',
    })
  }
}
```

```

    default: false,
}),
'config-dir': Flags.string({
  description: 'Override config directory',
  default: config.app.dataDir;
}),
};

async init(): Promise<void> {
  if (this.flags.debug) {
    this.logger.setLevel('debug');
  }
}
}
}

```

Jour 23-24: Core Commands Implementation

- [] **Tâche 6.4:** Implement init command

- **Owner:** Senior Dev
- **Deliverable:** src/commands/init.ts
- **Effort:** 4h
- **Functionality:**

```
// src/commands/init.ts
export default class Init extends BaseCommand {
  async run(): Promise<void> {
    // 1. Ask for workspace name
    // 2. Create directory structure
    // 3. Initialize SQLite database
    // 4. Create default config file
    // 5. Validate Gemini API key
    // 6. Test API connection
    this.log('✓ Workspace initialized at ./auto-gemini');
  }
}
```

- [] **Tâche 6.5:** Implement task run command

- **Owner:** Senior Dev
- **Deliverable:** src/commands/task/run.ts
- **Effort:** 6h
- **Functionality:**

```
// src/commands/task/run.ts
// Key features:
// - Load task from database
// - Execute with Gemini
// - Show live progress (spinner)
// - Stream output to terminal
// - Save result to session
```

Semaine 6 Closure

- [] **Tâche 6.6:** Test basic CLI workflow
 - **Owner:** Senior Dev
 - **Deliverable:** CLI test results
 - **Effort:** 2h
 - **Test:**
\$ npm run dev -- init --name test-workspace
\$ npm run dev -- task new "hello world"
\$ npm run dev -- task run <id>

Expected: Should complete without errors

Semaine 7 (17-23 février 2026)

Jour 25-26: TUI Components

- [] **Tâche 7.1:** Create TUI components (Ink.js or Blessed)
 - **Owner:** Senior Dev
 - **Deliverable:** TUI component library
 - **Effort:** 8h
 - **Components:**
// src/tui/components/
|—— Spinner.tsx # Loading spinner
|—— ProgressBar.tsx # Progress indication
|—— StatusBar.tsx # Bottom status bar
|—— TaskOutput.tsx # Streaming output display
|—— ErrorDisplay.tsx # Error messages with colors
|—— InteractivePrompt.tsx # User input
- [] **Tâche 7.2:** Implement chat interactive mode
 - **Owner:** Senior Dev
 - **Deliverable:** src/commands/chat.ts
 - **Effort:** 6h
 - **Features:**
 - [x] REPL interface (readline-based)
 - [x] Session persistence
 - [x] Streaming responses
 - [x] Command history
 - [x] Exit gracefully

Jour 27-28: Session Management UI

- [] **Tâche 7.3:** Implement session list and session show
 - **Owner:** Senior Dev
 - **Deliverable:** src/commands/session/*.ts
 - **Effort:** 5h
 - **Output format:**
\$ auto-gemini session list
ID | Created | Tasks | Last Modified
-----+--+

```
sess-001 | 2026-01-20 | 5 | 1 hour ago
sess-002 | 2026-01-15 | 12 | 2 days ago
$ auto-gemini session show sess-001
Session: sess-001
Created: 2026-01-20 14:30:00
Tasks executed: 5
Total tokens: 45,000
History:
[14:30] Task: "write hello world" → ✓
[14:32] Task: "optimize code" → ✓
```

...

- [] **Tâche 7.4:** Implement config commands
 - **Owner:** Senior Dev
 - **Deliverable:** src/commands/config.ts
 - **Effort:** 3h

Semaine 8 (24-2 mars 2026)

Jour 29-30: Feature Completeness

- [] **Tâche 8.1:** Implement remaining commands
 - **Owner:** Senior Dev
 - **Deliverable:** All commands from Phase 2 spec
 - **Effort:** 8h
 - **Commands to complete:**
 - task list, task delete
 - logs tail, logs grep
 - status --watch
 - All help pages
- [] **Tâche 8.2:** Test complete CLI flow
 - **Owner:** Senior Dev
 - **Deliverable:** Comprehensive CLI tests
 - **Effort:** 6h
 - **Test scenarios:**

Scenario 1: New user workflow

```
$ auto-gemini init
$ auto-gemini task new "write hello world"
$ auto-gemini task run <id>
$ auto-gemini session show <id>
```

Scenario 2: Chat mode

```
$ auto-gemini chat
Hello
What is my name?
exit
```

Scenario 3: Configuration

```
$ auto-gemini config set model gemini-2.0-pro  
$ auto-gemini config get model  
$ auto-gemini config show
```

Jour 31: Polish & Phase 2 Gate

- [] **Tâche 8.3:** Code cleanup & documentation
 - **Owner:** Senior Dev + Lead Arch
 - **Deliverable:** Clean codebase
 - **Effort:** 4h
 - **Checklist:**
 - [x] Remove console.log() (use logger)
 - [x] Add JSDoc comments
 - [x] CLI help text complete
 - [x] No broken links in docs
- [] **Tâche 8.4:** Phase 2 Gate Decision
 - **Owner:** Lead Arch
 - **Deliverable:** Go/No-Go for Phase 3
 - **Effort:** 2h
 - **Gate criteria:**

Phase 2 Gate Criteria

- ✓ Acceptance (GO to Phase 3) if:
 - [] All core commands working (init, task, chat, config)
 - [] CLI responsive and user-friendly
 - [] TUI components smooth (no lag)
 - [] Session persistence working
 - [] Help documentation complete
 - [] No critical bugs
- ✗ Re-planning (NO-GO) if:
 - [] TUI performance poor (>500ms latency)
 - [] Commands broken or inconsistent
 - [] Session data corrupt

Decision: GO NO-GO

PHASE 3: OPTIMISATION GEMINI (Semaines 9-12)

Objectif: Optimiser pour Gemini specifics (rate limiting, context, caching)

Livrable: Production-ready system, performance baseline, documentation

Succès: Passe load tests, handles constraints gracefully, monitoring active

Semaine 9 (3-9 mars 2026)

Jour 32-33: Rate Limiting System

- [] **Tâche 9.1:** Implement token-bucket rate limiter
 - **Owner:** Senior Dev
 - **Deliverable:** src/services/rate-limiterts
 - **Effort:** 6h
 - **Strategy:**

```
// Rate limit: 1000 requests / 24 hours
// Strategy: Token-bucket queue with SQLite persistence
interface RateLimitConfig {
  maxRequests: number; // 1000
  windowSeconds: number; // 86400 (24 hours)
  burstSize: number; // Allow 10 consecutive
}
class RateLimiter {
  async canMakeRequest(): Promise<boolean> {
    // Check if we have tokens available
    // If yes: consume token, return true
    // If no: queue request, return false
  }
  async queueRequest(task: Task): Promise<QueueID> {
    // Persist to DB with scheduled execution time
    // Emit event when task becomes executable
  }
}
```
- [] **Tâche 9.2:** Implement request queue & scheduler
 - **Owner:** Senior Dev
 - **Deliverable:** Queue management system
 - **Effort:** 6h
 - **Features:**
 - [x] Persistent queue (SQLite)
 - [x] Scheduled execution (cron-like)
 - [x] Priority levels (high, normal, low)
 - [x] Retry logic (exponential backoff)

Jour 34-35: Context Window Management

- [] **Tâche 9.3:** Implement context manager for 1M tokens
 - **Owner:** Senior Dev
 - **Deliverable:** src/services/context-managers
 - **Effort:** 8h
 - **Strategy:**

```
class ContextManager {
  maxTokens: number = 1_000_000;
  currentTokens: number = 0;
  async validatePrompt(prompt: string): Promise<ValidationResult> {
    const tokens = await this.countTokens(prompt);
    if (currentTokens + tokens > maxTokens) {
      return {
        valid: false,
    }
  }
}
```

- ```

availableTokens: maxTokens - currentTokens,
suggestedChunking: calculateChunks(prompt),
};

}

return { valid: true };
}

async addToContext(content: string): Promise<void> {
const tokens = await this.countTokens(content);
this.currentTokens += tokens;
}
}

```
- [ ] **Tâche 9.4:** Implement context window cleanup
    - **Owner:** Senior Dev
    - **Deliverable:** Sliding window & compression
    - **Effort:** 4h
    - **Strategies:**
      - Sliding window (remove oldest messages)
      - Compression (summarize old context)
      - Chunking (split large tasks)

## Semaine 10 (10-16 mars 2026)

### Jour 36-37: Performance Optimization

- [ ] **Tâche 10.1:** Optimize Gemini API calls
  - **Owner:** Senior Dev
  - **Deliverable:** Performance baseline & optimizations
  - **Effort:** 8h
  - **Optimizations:**

# Performance Optimizations

1. **Streaming Optimization**
    - Enable streaming for large outputs
    - Emit chunks as they arrive (don't buffer)
    - Reduces apparent latency
  2. **Caching Strategy**
    - Cache countTokens() results (same prompt)
    - Cache API responses (if safe)
    - Use Gemini Cache API (if available)
  3. **Batch Requests** (future)
    - If multiple tasks queued, batch API calls
    - Reduce API overhead
  4. **Model Selection**
    - Small tasks: gemini-2.0-flash (faster, cheaper)
    - Complex tasks: gemini-2.0-pro (more powerful)
    - Auto-detect based on complexity
- [ ] **Tâche 10.2:** Implement monitoring & metrics
    - **Owner:** Senior Dev
    - **Deliverable:** Metrics collection system
    - **Effort:** 5h

- Metrics to track:

## Key Metrics

- API Response Time (p50, p95, p99)
- Tokens Used / Task
- Queue Depth
- Error Rate by Type
- Session Duration
- Success Rate

Tools:

- Prometheus (optional, for advanced monitoring)
- Simple JSON logging for now

### Jour 38-39: Testing & Validation

- [ ] Tâche 10.3: Create load tests

- Owner: Senior Dev
- Deliverable: Load testing suite
- Effort: 6h
- Scenarios:
 

```
// Load test: Execute 100 tasks sequentially
// Measure: API latency, queue behavior, memory growth
async function loadTest() {
 for (let i = 0; i < 100; i++) {
 const task = task-${i};
 const start = Date.now();
 await agent.executeTask(task);
 const duration = Date.now() - start;
 console.log(Task ${i}: ${duration}ms);
 }
}
// Expected results:
// - p50: < 500ms
// - p95: < 2000ms
// - Memory: stable (no leaks)
// - Queue: processes at rate limit
```

- [ ] Tâche 10.4: Create stress tests

- Owner: Senior Dev
- Deliverable: Stress testing suite
- Effort: 4h
- Scenarios:
  - 1000+ token context
  - Very long prompts
  - Rapid-fire requests
  - Network failures/timeouts

## Semaine 11 (17-23 mars 2026)

### Jour 40-41: Error Handling & Resilience

- [ ] **Tâche 11.1:** Comprehensive error handling
  - **Owner:** Senior Dev
  - **Deliverable:** Error recovery strategies for all Gemini errors
  - **Effort:** 6h
  - **Error types handled:**
    - [x] Rate limiting (RESOURCE\_EXHAUSTED)
    - [x] Auth failures (PERMISSION\_DENIED)
    - [x] Content safety (INVALID\_ARGUMENT)
    - [x] Context window overflow
    - [x] Network timeouts
    - [x] Malformed responses
- [ ] **Tâche 11.2:** Implement circuit breaker
  - **Owner:** Senior Dev
  - **Deliverable:** Circuit breaker pattern implementation
  - **Effort:** 4h
  - **States:**
    - CLOSED: Normal operation
    - OPEN: API failing, reject requests
    - HALF\_OPEN: Testing if API recovered

### Jour 42-43: Documentation & Examples

- [ ] **Tâche 11.3:** Write comprehensive documentation
  - **Owner:** Senior Dev + Lead Arch
  - **Deliverable:** /docs/USER\_GUIDE.md, /docs/API.md
  - **Effort:** 8h
  - **Content:**

# User Guide

## Installation

```
npm install -g auto-gemini-cli
```

or

```
npx auto-gemini-cli init
```

## Quick Start

```
$ auto-gemini init
$ auto-gemini task new "write hello world"
$ auto-gemini task run <id>
```

# Configuration

## API Key

```
export GEMINI_API_KEY="..."
auto-gemini config set apiKey $GEMINI_API_KEY
```

## Rate Limiting

- Free tier: 1000 requests/24h
- Monitor queue: auto-gemini status --watch
- Scheduled execution: Automatic

# Troubleshooting

## Rate Limit Exceeded

- Tasks queued automatically
- Check status: auto-gemini status
- Manually clear: auto-gemini session clear <id>

## Context Window Full

- Auto-chunking enabled
  - Or: Reduce task scope
  - Or: Start new session
- [ ] **Tâche 11.4:** Create example workflows
    - **Owner:** Senior Dev
    - **Deliverable:** Example scripts & documentation
    - **Effort:** 4h
    - **Examples:**
      - Code generation workflow
      - Documentation writing
      - Bug analysis workflow

## Semaine 12 (24-30 mars 2026)

### Jour 44-45: Integration Testing

- [ ] **Tâche 12.1:** Run comprehensive integration tests
  - **Owner:** Senior Dev
  - **Deliverable:** Integration test suite results
  - **Effort:** 8h
  - **Test coverage:**

# Integration Tests

- [ ] Rate limiter queues requests correctly
- [ ] Context manager prevents overflow
- [ ] Error recovery works for each error type
- [ ] Streaming continues after network hiccup
- [ ] Sessions persist across restarts
- [ ] Parallel tasks don't corrupt state
- [ ] Memory stable over 24h simulation

- [ ] **Tâche 12.2:** Performance benchmarking
  - **Owner:** Senior Dev
  - **Deliverable:** Performance report
  - **Effort:** 4h
  - **Baseline metrics:**
    - Task Execution Time:
      - Simple (< 100 tokens): 200-500ms
      - Medium (100-1000 tokens): 500-2000ms
      - Large (1000+ tokens): 2000-5000ms
    - Memory Usage:
      - Idle: < 50MB
      - Single session: < 100MB
      - Multi-session (10): < 200MB
    - API Overhead:
      - countTokens() call: ~50ms
      - generateContent() call: ~200-1000ms (varies)

#### Jour 46-47: Phase 3 Gate

- [ ] **Tâche 12.3:** Production readiness review
  - **Owner:** Lead Arch + Senior Dev
  - **Deliverable:** Production readiness checklist
  - **Effort:** 6h
  - **Checklist:**

## Production Readiness Checklist

### Code Quality

- [ ] No console.log() statements
- [ ] All errors have appropriate recovery
- [ ] Logging is comprehensive
- [ ] Type safety: no any types
- [ ] Performance baseline established
- [ ] Memory leaks tested & verified none

### Operations

- [ ] Health check endpoint works
- [ ] Logging to file for analysis
- [ ] Error tracking integrated (Sentry)
- [ ] Metrics exported (Prometheus optional)
- [ ] Documentation complete & accurate

### Security

- [ ] API key never logged
- [ ] Secrets not in repository
- [ ] SQLite database encrypted (optional)
- [ ] Rate limiting prevents abuse
- [ ] Input validation comprehensive

## Testing

- [ ] Unit tests: > 80% coverage
- [ ] Integration tests: all critical paths
- [ ] Load tests: < 2s p95 latency
- [ ] Stress tests: handles edge cases
- [ ] E2E tests: against real Gemini API

## Deployment

- [ ] Build process automated
  - [ ] Docker image created (optional)
  - [ ] Executable binary tested
  - [ ] Installation documented
  - [ ] Upgrade path documented
- [ ] **Tâche 12.4:** Phase 3 Gate Decision
    - **Owner:** Lead Arch
    - **Deliverable:** Go/No-Go for Phase 4
    - **Effort:** 2h
    - **Gate criteria:**

## Phase 3 Gate Criteria

- ✓ Acceptance (GO to Phase 4) if:
  - [ ] Rate limiting working correctly
  - [ ] Context manager prevents overflows
  - [ ] Performance meets baseline
  - [ ] Error handling comprehensive
  - [ ] Load tests pass (p95 < 2s)
  - [ ] No known bugs
  - [ ] Production readiness > 90%
- ✗ Re-planning (NO-GO) if:
  - [ ] Performance far below baseline
  - [ ] Rate limiting buggy
  - [ ] Memory leaks detected

Decision:  GO  NO-GO

---

## PHASE 4: INTÉGRATION & RELEASE (Semaines 13-14)

**Objectif:** Tester e2e, finaliser documentation, déployer v1.0

**Livrable:** v1.0 release, documentation complète, support ready

**Succès:** v1.0 dans npm registry, utilisable en production

Semaine 13 (31 mars - 6 avril 2026)

## Jour 48-49: Final Testing

- [ ] **Tâche 13.1:** Execute comprehensive E2E tests
  - **Owner:** Senior Dev
  - **Deliverable:** E2E test results (real Gemini API)
  - **Effort:** 8h
  - **Test scenarios:**

### Scenario 1: Fresh install

```
npm install -g auto-gemini-cli
auto-gemini init
auto-gemini config set apiKey $GEMINI_API_KEY
auto-gemini task new "write a fibonacci function"
auto-gemini task run <id>
```

**Expected: Task completes successfully**

**Result stored in session**

### Scenario 2: Long-running task

```
auto-gemini task new "analyze this large repository" # 500K+ tokens
auto-gemini task run <id> --watch
```

**Expected: Auto-chunks, shows progress, completes**

### Scenario 3: Rate limiting

```
for i in {1..20}; do
 auto-gemini task new "task i" auto-gemini task run task-i --no-wait
done
auto-gemini status --watch
```

**Expected: First 10 run, rest queued**

**Queue processes as rate limit allows**

# Scenario 4: Error recovery

(Manually trigger network failure)

**Expected: Circuit breaker engages, retries work**

- [ ] **Tâche 13.2:** User acceptance testing (if stakeholders available)
  - **Owner:** Desjardins stakeholders + Lead Arch
  - **Deliverable:** Feedback & sign-off
  - **Effort:** 4h
  - **Feedback areas:**
    - CLI UX (is it intuitive?)
    - Performance (is it fast enough?)
    - Documentation (is it clear?)
    - Reliability (does it work consistently?)

Jour 50-51: Documentation & Examples

- [ ] **Tâche 13.3:** Finalize documentation
  - **Owner:** Senior Dev + Lead Arch
  - **Deliverable:** Complete documentation suite
  - **Effort:** 8h
  - **Documents to create/update:**
    - [README.md](#) (overview, quick start)
    - [INSTALLATION.md](#) (detailed install)
    - [USER\\_GUIDE.md](#) (step-by-step)
    - [API.md](#) (command reference)
    - [TROUBLESHOOTING.md](#) (common issues)
    - [ARCHITECTURE.md](#) (technical details)
    - [MIGRATION.md](#) (from Auto-Claude)
    - [CONTRIBUTING.md](#) (for future contributors)
- [ ] **Tâche 13.4:** Create changelog & release notes
  - **Owner:** Lead Arch
  - **Deliverable:** [CHANGELOG.md](#), release notes
  - **Effort:** 2h
  - **Content:**

## CHANGELOG

v1.0.0 - 2026-05-01

## Major Features

- ★★ Gemini CLI: Full adaptation of Auto-Claude for Gemini API
- ★★ Rate limiting: Automatic queue management (1000 req/24h)
- ★★ Context management: Handle up to 1M token contexts
- ★★ Interactive chat: REPL mode with session persistence
- ★★ Performance: Streaming responses, fast token counting

## Breaking Changes

- Electron UI removed (CLI only)
- Config file format slightly different
- API key env var: GEMINI\_API\_KEY (was CLAUDE\_API\_KEY)

## Migration

- See [MIGRATION.md](#) for moving from Auto-Claude

## Known Limitations

- Free tier: 1000 requests/24h (Google limit)
- No multi-model support yet
- No caching API integration yet

## Contributors

- André-Guy Bruneau (Lead Architect)
- [Team members]

## Semaine 14 (7-13 avril 2026)

### Jour 52-53: Build & Publishing

- [ ] **Tâche 14.1:** Create build artifacts
  - **Owner:** Lead Arch
  - **Deliverable:** Packaged binaries & npm package
  - **Effort:** 4h
  - **Artifacts:**

## Build executable

npm run build

## Package for different platforms (optional)

npm run package:macos  
npm run package:linux  
npm run package:windows

## Output:

dist/

  └── auto-gemini-cli-1.0.0.tar.gz

  └── auto-gemini-cli-1.0.0.exe

    └── auto-gemini-cli-1.0.0.dmg

- [ ] **Tâche 14.2:** Publish to npm registry
  - **Owner:** Lead Arch
  - **Deliverable:** npm package published
  - **Effort:** 2h
  - **Steps:**

## Verify package.json

npm version major # This is v1.0.0

## Test publish (optional)

npm publish --dry-run

## Publish to npm

npm publish

## Verify

npm search auto-gemini-cli

npm view auto-gemini-cli@1.0.0

- [ ] **Tâche 14.3:** Create GitHub release
  - **Owner:** Lead Arch
  - **Deliverable:** GitHub release page
  - **Effort:** 1h
  - **Content:**
    - Release notes (copy from CHANGELOG)
    - Binaries attached
    - Installation instructions
    - Known issues

## Jour 54-55: Release & Support

- [ ] **Tâche 14.4:** Announce release
  - **Owner:** Lead Arch
  - **Deliverable:** Release announcement
  - **Effort:** 2h
  - **Channels:**
    - GitHub releases
    - Email to stakeholders
    - Internal documentation
    - (Optional: blog post, Twitter, etc.)
- [ ] **Tâche 14.5:** Setup support & issue tracking
  - **Owner:** Lead Arch
  - **Deliverable:** Support guidelines, issue templates
  - **Effort:** 2h
  - **Checklist:**

## Support Setup

- [ ] GitHub issues enabled
- [ ] Issue templates created (.github/ISSUE\_TEMPLATE/)
- [ ] Discussions enabled
- [ ] Slack channel #auto-gemini-cli-support
- [ ] Email support alias set up
- [ ] Response time SLAs defined

### Issue Templates:

- Bug Report
- Feature Request
- Question
- Documentation Issue

## Jour 56: Final Validation & Gate

- [ ] **Tâche 14.6:** v1.0 Sign-off
  - **Owner:** Lead Arch + Desjardins Leadership
  - **Deliverable:** Official sign-off
  - **Effort:** 2h
  - **Final checklist:**

## v1.0 Release Sign-Off

### Technical

- [ ] All tests passing (unit, integration, E2E)
- [ ] Performance baseline met
- [ ] No critical bugs
- [ ] Documentation complete
- [ ] Code reviewed & approved

## Operational

- [ ] npm package published
- [ ] GitHub release created
- [ ] Support process ready
- [ ] Monitoring active

## Stakeholder

- [ ] Desjardins approval received
- [ ] Budget tracking complete
- [ ] Timeline met (or documented variance)
- [ ] Deliverables accepted

**Sign-off:**  APPROVED FOR PRODUCTION

Date: \_\_\_\_\_

By: André-Guy Bruneau (Lead Architect)

For: Desjardins

- [ ] **Tâche 14.7:** Post-Release Planning
  - **Owner:** Lead Arch
  - **Deliverable:** Roadmap for future versions
  - **Effort:** 2h
  - **Future work:**

## Post-v1.0 Roadmap

### v1.1 (2-3 months after v1.0)

- [ ] Gemini Cache API integration
- [ ] Multi-model support (Claude + Gemini)
- [ ] Web UI (optional)
- [ ] Plugin system

### v2.0 (6+ months after v1.0)

- [ ] Distributed execution (multi-machine)
- [ ] Advanced scheduling
- [ ] Workflow orchestration
- [ ] Team collaboration features

## Community

- [ ] Contribute guidelines
- [ ] Plugin template
- [ ] Community examples repo

---

## DÉPENDANCES & BLOCKERS

## Dépendances Critiques

Phase 0:

- └→ Approval budgétaire Desjardins (CRITICAL)
- └→ Phase 1 begins

Phase 1:

- └→ Gemini API stable & available (CRITICAL)
- └→ Codebase analysis complete
- └→ Phase 2 begins

Phase 2:

- └→ Phase 1 complete (at least MVP working)
- └→ Oclif framework understood
- └→ Phase 3 begins

Phase 3:

- └→ Phase 2 complete (CLI working)
- └→ Load testing tools available
- └→ Phase 4 begins

Phase 4:

- └→ Phase 3 complete (optimization done)
- └→ npm account available
- └→ Release to npm

## Potential Blockers & Mitigation

| Blocker                                        | Impact          | Probability | Mitigation                        |
|------------------------------------------------|-----------------|-------------|-----------------------------------|
| <b>Gemini API breaking change</b>              | 2-3 weeks delay | MEDIUM      | Monitor API status, use adapters  |
| <b>Codebase more coupled than expected</b>     | 1-2 weeks delay | MEDIUM      | Design refactoring upfront        |
| <b>Team member unavailable</b>                 | Timeline slip   | LOW         | Cross-train both team members     |
| <b>Performance unacceptable</b>                | 1 week delay    | LOW         | Optimize streaming, caching       |
| <b>Rate limiting impossible to work around</b> | Pivot needed    | LOW         | Fall back to Scenario 3 (Claude)  |
| <b>Gemini free tier discontinued</b>           | Costs increase  | LOW         | Switch to paid tier or Scenario 3 |

### Risk Mitigation Actions

- **Weekly risk review:** Every Monday, assess blockers
- **Slack #auto-gemini-risks:** Real-time communication
- **Decision framework:** If blocker > 2 days impact, escalate immediately
- **Buffer time:** 2-3 weeks reserved for unexpected delays

---

## MÉTRIQUES DE SUCCÈS

Phase-by-Phase Success Metrics

| Phase | Key Metrics                | Target | Measurement            |
|-------|----------------------------|--------|------------------------|
| 0     | Approval %                 | 100%   | Budget signed          |
|       | Codebase understood        | Yes    | ADRs written           |
|       | Team ready                 | Yes    | Env setup complete     |
| 1     | API swap working           | Yes    | E2E test passes        |
|       | Tests passing              | 70%+   | Test report            |
|       | Token counter validated    | Yes    | countTokens tests      |
| 2     | Commands functional        | 100%   | CLI tests pass         |
|       | TUI responsive             | <500ms | Latency test           |
|       | User can complete workflow | Yes    | User test successful   |
| 3     | Rate limiting works        | Yes    | Queue test passes      |
|       | Context management works   | Yes    | 1M token test          |
|       | Performance baseline met   | Yes    | Benchmarks pass        |
| 4     | v1.0 released to npm       | Yes    | npm publish successful |
|       | Documentation complete     | Yes    | All docs reviewed      |
|       | Zero critical bugs         | Yes    | Bug report review      |

## Overall Project Success Metrics

Timeline: 14 weeks (3.5 months) ✓ vs original 19 weeks

Cost: \$13-34K ✓ vs original \$100-150K

Features: 100% of Phase 1-4 scope ✓

Quality: >80% test coverage ✓

Performance: p95 < 2s ✓

Uptime (day 1 of release): >95% ✓

User satisfaction: >4/5 ✓

Zero production hotfixes needed ✓

# GESTION DES RISQUES

## Risk Register

### Risk #1: Gemini API Incompatibility

- **Description:** Critical incompatibility between Claude and Gemini APIs discovered during Phase 1
- **Probability:** MEDIUM (50%)
- **Impact:** HIGH (2-3 weeks delay)
- **Mitigation:** Adapter pattern designed to isolate API differences
- **Owner:** Senior Dev
- **Action:** Daily API testing, early detection

### Risk #2: Performance Unacceptable

- **Description:** Gemini API slower than Claude, impacting user experience
- **Probability:** LOW (20%)
- **Impact:** MEDIUM (1 week delay)
- **Mitigation:** Implement streaming, caching, optimize before Phase 4
- **Owner:** Senior Dev
- **Action:** Performance profiling in Phase 3

### Risk #3: Rate Limiting Not Workable

- **Description:** 1000 req/24h limit too restrictive for use cases
- **Probability:** LOW (15%)
- **Impact:** HIGH (complete pivot)
- **Mitigation:** Pivot to Scenario 3 (Claude Code) as backup
- **Owner:** Lead Arch
- **Action:** Have Scenario 3 plan ready

### Risk #4: Team Member Unavailable

- **Description:** One team member becomes unavailable mid-project
- **Probability:** LOW (10%)
- **Impact:** MEDIUM (1-2 weeks delay)
- **Mitigation:** Cross-train both team members early
- **Owner:** Lead Arch
- **Action:** Shared knowledge sessions Week 1-2

### Risk #5: Scope Creep

- **Description:** Additional features requested mid-project
- **Probability:** MEDIUM (60%)
- **Impact:** MEDIUM (timeline slip)
- **Mitigation:** Strict change control, defer to v1.1
- **Owner:** Lead Arch
- **Action:** Weekly stakeholder alignment

## Risk Monitoring

- **Weekly risk review:** Every Monday in standup
  - **Escalation path:** If risk impacts > 2 days, escalate to Lead Arch
  - **Risk board:** GitHub project "Risks" tab
  - **Contingency time:** 2-3 weeks buffer built into timeline
- 

# BUDGET DÉTAILLÉ

## Cost Breakdown

### PHASE 0: Préparation (2 weeks)

|—— Lead Architect (1 FTE):  $1.0 \times 2 \times \$200/\text{hr} = \$1,600$   
|—— Senior Dev (0.5 FTE):  $0.5 \times 2 \times \$200/\text{hr} = \$800$   
|—— Subtotal: = \$2,400

### PHASE 1: Swap API (3 weeks)

|—— Lead Architect (1 FTE):  $1.0 \times 3 \times \$200/\text{hr} = \$2,400$   
|—— Senior Dev (0.5 FTE):  $0.5 \times 3 \times \$200/\text{hr} = \$1,200$   
|—— Testing tools (optional): = \$500  
|—— Subtotal: = \$4,100

### PHASE 2: UI Migration (3 weeks)

|—— Lead Architect (0.5 FTE):  $0.5 \times 3 \times \$200/\text{hr} = \$1,200$   
|—— Senior Dev (1 FTE):  $1.0 \times 3 \times \$200/\text{hr} = \$2,400$   
|—— Design review: = \$300  
|—— Subtotal: = \$3,900

### PHASE 3: Optimization (4 weeks)

|—— Lead Architect (0.5 FTE):  $0.5 \times 4 \times \$200/\text{hr} = \$1,600$   
|—— Senior Dev (1 FTE):  $1.0 \times 4 \times \$200/\text{hr} = \$3,200$   
|—— Load testing tools: = \$500  
|—— Monitoring (Sentry free): = \$0  
|—— Subtotal: = \$5,300

### PHASE 4: Release (2 weeks)

|—— Lead Architect (1 FTE):  $1.0 \times 2 \times \$200/\text{hr} = \$1,600$   
|—— Senior Dev (0.5 FTE):  $0.5 \times 2 \times \$200/\text{hr} = \$800$   
|—— Documentation review: = \$300  
|—— Subtotal: = \$2,700

## INFRASTRUCTURE & TOOLS:

|—— GitHub (free): = \$0  
|—— npm (free): = \$0  
|—— Google Gemini (free tier): = \$0  
|—— Sentry (free): = \$0  
|—— Slack (existing): = \$0  
|—— Subtotal: = \$0

---

TOTAL PROJECT COST: \$18,400

Contingency (10% for unknowns): \$ 1,840

---

TOTAL BUDGET (with contingency): \$20,240

APPROVED BUDGET RANGE: \$13-34K

STATUS: ✓ WITHIN RANGE

## Budget Allocation by Cost Center

Labor Costs (FTE Hours):

- └── Lead Architect:  $1.0 \times 14 \text{ weeks} \times 40\text{h/week} \times \$200/\text{hr} = \$11,200$
- └── Senior Dev:  $0.5 \times 14 \text{ weeks} \times 40\text{h/week} \times \$200/\text{hr} = \$5,600$
- └── Total Labor: = \$16,800

Tools & Services: = \$1,300

- └── Testing tools: \$500
- └── Monitoring: \$300 (Sentry optional)
- └── Design review: \$300
- └── Documentation: \$200

Contingency (10%): = \$1,840

TOTAL: = \$20,240

## Cost Savings Opportunities

- **Reuse Auto-Claude code:** Saves ~30% development (vs Scenario 1)
- **Leverage Gemini free tier:** No API costs (1000 req/24h)
- **Use free tools:** GitHub, npm, Sentry free tier
- **Cross-training:** Avoids hiring additional staff

## Cost Tracking

Weekly cost tracking:

- Team hours logged (time tracking)
- Tool subscriptions tracked
- Contingency reserve managed
- Budget variance reported

Monthly stakeholder report:

- Actual vs planned spending
  - Forecast to completion
  - Risk adjustments
-

# CHECKLIST FINALE

## Pre-Launch (Semaine 13)

- [ ] All phases completed
- [ ] All tests passing (>90%)
- [ ] Documentation reviewed
- [ ] Performance baseline met
- [ ] Security audit passed
- [ ] Stakeholder sign-off obtained

## Launch (Semaine 14)

- [ ] npm package published
- [ ] GitHub release created
- [ ] Announcement sent
- [ ] Support team ready
- [ ] Monitoring active
- [ ] Rollback plan documented

## Post-Launch (First 2 weeks)

- [ ] Monitor error rates (target: <1%)
- [ ] Track user feedback
- [ ] Respond to support requests (<4h)
- [ ] Document lessons learned
- [ ] Plan v1.1 features

---

# CONCLUSION

## Planification complète pour Scénario 2: Adaptation Auto-Claude vers Auto-Gemini-CLI

- **Durée totale:** 14 semaines (3.5 mois), livraison ~1er mai 2026
- **Effort:** 1.5 FTE (Lead Architect + Senior Developer)
- **Budget:** \$13-34K (plus ou moins contingency)
- **Approche:** Itérative par phases, avec gates de décision claire
- **Succès défini:** v1.0 production-ready, documenté, testé, relâché

## Prochaines étapes:

1. Obtenir approbation budgétaire formelle de Desjardins
2. Assembler équipe (Lead Arch + Senior Dev)
3. Commencer Phase 0 semaine de 6 janvier 2026
4. Hold bi-weekly stakeholder reviews

**Contact & Questions:** André-Guy Bruneau (agbruneau)